

# CodeAgent Vulnerability Scanner — Complete API & Operations Guide (FastAPI)

Base URL (dev): `http://localhost:8080`

This document is the authoritative reference for the CodeAgent Vulnerability Scanner service. It covers **how the API works end-to-end**, all endpoints and models, job lifecycle, rate limits, auth, error codes, webhooks/SSE, analyzer configuration, and implementation notes. It is designed so you can hand it to frontend devs, partner integrators, or your future self.

---

## 1) Overview & Architecture

**What it does:** Accepts a GitHub repo URL or a ZIP upload → clones/extracts → runs analyzers (Semgrep, Bandit, dependency audit) → produces a **normalized JSON report** grouped by file and severity → exposes it via `/reports/{job_id}`. Optional **async mode**, **SSE progress**, and **webhooks**. Emits an internal `report.created` event for the agentic layer (autofix/test/PR creation).

**Components - FastAPI app** (`/analyze`, `/reports`, `/jobs`, `/events`, `/webhooks`, `/config`). - **Ingestion:** safe clone/extract into `storage/workspace/{job_id}` with ignores. - **Analyzers:** plugin runners (Semgrep, Bandit, pip-audit), concurrent execution with timeouts. - **Orchestrator:** merges findings → normalized report → writes to `storage/reports/{job_id}.json`. - **Events:** optional webhooks (HTTP POST) and in-process bus for agentic subscribers.

---

## 2) Authentication, Headers, Formats

### 2.1 Auth

- **Recommended (prod):** `Authorization: Bearer <API_KEY>` (HMAC/DB check server-side)
- **Dev:** unauthenticated

### 2.2 Global Headers

- `Idempotency-Key: <uuid>` — prevents duplicate scans on retries
- `X-Client-Id: <string>` — (optional) tenant/tracking
- `X-Request-Id` — returned on each response; you can also supply one

### 2.3 Content Types

- Requests: `multipart/form-data` (uploads), `application/json` (control/config)
- Responses: `application/json` (except SSE stream)

## 2.4 Versions

- **Accept:** `application/vnd.codeagent.v1+json` (optional vendor versioning)
  - **Minor, backward-compatible changes bump** `version` field in `/health` and OpenAPI
- 

## 3) Limits, Quotas, Timeouts

- **Upload cap:** 50 MB (413 on exceed)
  - **Per job limits:** max files (e.g., 10,000), per-tool timeout (default 600s), overall wall-clock (e.g., 900s)
  - **Rate limit:** 60 req/min per API key; max 2 concurrent `running` jobs per key
  - **Ignore lists:** `.git`, `node_modules`, `venv`, `.venv`, `dist`, `build`, `__pycache__`, large binaries > 20MB
  - **URL allow-list:** `https://github.com/*` initially; can be extended via config
- 

## 4) Data Model

### 4.1 Severity Enum

`critical | high | medium | low`

### 4.2 Issue (normalized finding)

```
{
  "tool": "bandit",
  "type": "B608",
  "message": "Possible SQL injection via string building",
  "severity": "high",
  "file": "src/app.py",
  "line": 42,
  "rule_id": "B608",
  "suggestion": "Use parameterized queries"
}
```

### 4.3 FileIssues

```
{ "path": "src/app.py", "issues": [Issue, ...] }
```

### 4.4 Report

```
{
  "job_id": "4b3a...",

```

```

"meta": {
  "tools": ["semgrep", "bandit", "pip-audit"],
  "repo": {"source": "github|zip", "url": "https://github.com/org/repo",
"ref": "main", "commit": null},
  "generated_at": "2025-10-17T14:30:00Z",
  "duration_ms": 23456,
  "labels": ["customer:acme", "case:123"]
},
"summary": {"critical": 1, "high": 3, "medium": 7, "low": 4},
"files": [FileIssues, ...]
}

```

## 4.5 Error

```

{
  "error": {
    "code": "INVALID_INPUT|NOT_FOUND|TIMEOUT|INTERNAL|RATE_LIMIT|
PAYLOAD_TOO_LARGE|UNAUTHORIZED|FORBIDDEN",
    "message": "Human-readable",
    "details": {"field": "github_url"}
  }
}

```

## 5) Job Lifecycle

queued → running (clone → analyze:N tools → merge → write) → completed | failed | canceled | expired

- **Creation:** /analyze or /analyze-async returns job\_id
- **Progress:** /jobs/{job\_id} or SSE /events/{job\_id}
- **Result:** /reports/{job\_id} once completed
- **Cancel:** DELETE /jobs/{job\_id} when queued|running

## 6) Endpoints

### 6.1 Health & Metadata

GET /health → 200 { "status": "ok", "version": "0.1.0" }

GET /tools → 200 { "available": ["bandit", "semgrep", "dep"], "default": ["bandit", "semgrep", "dep"], "versions": {"bandit": "1.7.9", "semgrep": "1.81.0"} }

## 6.2 Submit Scan (Sync / Auto-Async)

**POST** `/analyze` — may reply `200` (small jobs) or `202` (accepted) based on size/timeout hints.

**Headers:** `Content-Type: multipart/form-data`, optional `Idempotency-Key`

**Form fields** - `github_url` (*string, optional*) — `https://github.com/org/repo` - `ref` (*string, optional*) — branch/tag - `commit` (*string, optional*) — SHA to checkout - `file` (*file, optional*) — ZIP archive - `include` (*string, optional*) — comma CSV of globs (`src/**/*.py`) - `exclude` (*string, optional*) — comma CSV of globs (`tests/**/*.py`) - `analyzers` (*string, optional*) — CSV: `bandit,semgrep,dep` - `timeout_sec` (*int, optional*) — overrides default - `labels` (*string, optional*) — CSV labels (`customer:acme,case:123`)

**Responses** - `200 OK` — `{ "job_id": "...", "summary": { ... } }` - `202 Accepted` — `{ "job_id": "...", "status": "running" }` - `400/401/403/413/429/500` — *Error*

**Validation rules** - Require **either** `github_url` **or** `file` (not both) - Validate `github_url` host and scheme; enforce allow-list - Enforce upload size cap and filename checks

---

## 6.3 Submit Scan (Async)

**POST** `/analyze-async`

**Request** — same fields as `/analyze`.

**202 Accepted**

```
{ "job_id": "uuid", "status": "queued" }
```

Background task performs clone/extract → analyze → write report → emit `report.created` and deliver webhooks.

---

## 6.4 Jobs

**GET** `/jobs/{job_id}` → status & timestamps

```
{
  "job_id": "...",
  "status": "queued|running|completed|failed|canceled|expired",
  "progress": { "phase": "clone|analyze:semgrep|analyze:bandit|merge|write",
  "percent": 65},
```

```
"submitted_at": "...",
"started_at": "...",
"finished_at": "...",
"error": null
}
```

**DELETE** `/jobs/{job_id}` — cancel - 202 { "job\_id": "...", "status": "canceling" } - 409 if not cancellable

**POST** `/jobs/{job_id}/rerun` — requeue with same inputs - 202 { "job\_id": "...", "status": "queued" }

## 6.5 Reports

**GET** `/reports/{job_id}` — full report - 200 Report - 404 Not found (not finished/doesn't exist)

**GET** `/reports` — paginate/filter **Query:** `page`, `limit` ( $\leq 100$ ), `severity=high,critical`, `tool=semgrep`, `repo=https://github.com/org/repo`, `since`, `until`, `label`

```
{
  "items": [Report, ...],
  "page": 1,
  "limit": 20,
  "total": 57
}
```

**GET** `/reports/{job_id}/summary` — light summary

```
{ "job_id": "...", "summary": {"critical":0,"high":2,"medium":7,"low":5} }
```

## 6.6 Live Progress (SSE)

**GET** `/events/{job_id}` - Header: `Accept: text/event-stream` - Server emits events:

```
:event: progress
:data: {"phase":"semgrep","percent":32}

:event: finished
:data: {"job_id":"...","status":"completed"}
```

---

## 6.7 Webhooks

**POST** /webhooks/register — register a URL to receive report.created

```
{ "url": "https://example.com/hooks", "events": ["report.created"], "secret": "optional" }
```

- 201 { "id": "wh\_123" }

**Delivery** - Method: POST - Headers: X-Event: report.created, X-Signature: sha256=... (HMAC over body) - Body:

```
{ "job_id": "...", "repo": {"url":"..."}, "summary": {...}, "report_url": "/reports/..." }
```

- Retries with exponential backoff on 5xx/timeouts; disable after N attempts

**DELETE** /webhooks/{id} — unregister

---

## 6.8 Configuration

**GET** /config/analyzers

```
{ "defaults": ["bandit","semgrep","dep"], "rulesets": {"semgrep":["p/owasp-top-ten","p/secrets"], "bandit": []}, "allow_list": ["https://github.com/"] }
```

**PATCH** /config/analyzers

```
{ "defaults": ["semgrep","dep"], "rulesets": {"semgrep":["p/owasp-top-ten","p/secrets"]}, "allow_list": ["https://github.com/", "https://gitlab.com/"] }
```

- 200 { "ok": true }

---

## 7) How Scanning Works (under the hood)

### 1. Ingestion

2. Validate inputs → create job\_id → make storage/workspace/{job\_id}
3. git clone --depth=1 {url} (or extract ZIP)

4. Remove ignored folders; optionally apply `include` / `exclude` globs
5. **Analyzer selection**
6. Parse `analyzers` CSV or use defaults
7. Enable Bandit only if `.py` files exist; Dep audit only if `requirements.txt` exists (JS/npm audit/OSV later)
8. **Execution**
9. Run in parallel (ThreadPoolExecutor), each with a per-tool timeout
10. Capture JSON output (Semgrep/Bandit/pip-audit); map to normalized Issue
11. **Merge**
12. Compute severity summary; group by file; attach metadata (tools, repo, duration)
13. **Persist & notify**
14. Write `storage/reports/{job_id}.json`
15. Update job state to `completed` or `failed`
16. Emit `report.created` (events + webhooks)

**Security note:** The service never executes repo code. It only runs static analyzers over files. All paths are confined under the job workspace.

---

## 8) Error Handling & Status Codes

- `200 OK` — success (sync)
- `202 Accepted` — queued/running (async or auto-async)
- `400 INVALID_INPUT` — missing/invalid fields; both `github_url` and `file` sent; bad URL host
- `401 UNAUTHORIZED` — missing/invalid API key (prod)
- `403 FORBIDDEN` — plan/tenant restrictions
- `404 NOT_FOUND` — unknown `job_id` /report
- `409 CONFLICT` — cannot cancel/rerun
- `413 PAYLOAD_TOO_LARGE` — upload exceeds limit
- `429 RATE_LIMIT` — throttled; retry after `Retry-After`
- `500 INTERNAL` — unexpected server error
- `504 TIMEOUT` — per-tool or overall timeout exceeded

Each error uses the **Error** shape from §4.5 and returns an `X-Request-Id` header.

---

## 9) OpenAPI (runnable excerpt)

```
openapi: 3.0.3
info:
  title: CodeAgent Scanner API
  version: 0.1.0
servers:
  - url: http://localhost:8080
paths:
  /health:
```

```

get:
  summary: Health check
  responses:
    '200': { description: OK }
/analyze:
  post:
    summary: Submit a scan (may return 200 or 202)
    requestBody:
      required: true
      content:
        multipart/form-data:
          schema:
            type: object
            properties:
              github_url: { type: string, format: uri }
              ref: { type: string }
              commit: { type: string }
              file: { type: string, format: binary }
              include: { type: string }
              exclude: { type: string }
              analyzers: { type: string }
              timeout_sec: { type: integer }
              labels: { type: string }
    responses:
      '200': { description: OK }
      '202': { description: Accepted }
      '400': { description: Bad Request }
      '413': { description: Payload Too Large }
      '500': { description: Internal Error }
/reports/{job_id}:
  get:
    summary: Get full report
    parameters:
      - in: path
        name: job_id
        required: true
        schema: { type: string }
    responses:
      '200': { description: OK }
      '404': { description: Not Found }

```

Full OpenAPI can be generated directly by FastAPI from the routers; this excerpt anchors required fields.



## 10) Usage Examples

### GitHub URL (auto-async on big repos)

```
curl -F github_url=https://github.com/pallets/flask
-F analyzers=bandit,semgrep,dep
http://localhost:8080/analyze
```

### ZIP Upload (explicit async)

```
curl -F file=@/path/to/project.zip
-F include="src/**/*.py"
http://localhost:8080/analyze-async
```

### Get Status & Report

```
curl http://localhost:8080/jobs/<job_id> | jq
curl http://localhost:8080/reports/<job_id> | jq
```

### Register Webhook

```
curl -X POST http://localhost:8080/webhooks/register
-H 'Content-Type: application/json'
-d '{"url":"https://example.com/hooks","events":
["report.created"],"secret":"xyz"}
```

---

## 11) Implementation Notes (for engineers)

- **Routers:** `api/routers/analyze.py` implements `/analyze`, `/analyze-async`, `/jobs/*`, `/reports/*`, `/events/*`.
  - **BackgroundTasks:** used to offload large jobs; swap to Celery/RQ later for retries.
  - **ThreadPool:** used per job to run analyzers in parallel; cap with env vars.
  - **Storage layout:**
    - `storage/workspace/{job_id}` — cloned/extracted repo
    - `storage/reports/{job_id}.json` — final report
    - `storage/logs/{job_id}.json` — status/progress, errors
  - **Security:** never execute repo code; sanitize paths; enforce allow-lists; set subprocess timeouts.
  - **Observability:** log JSON with `job_id` correlation; expose `X-Request-Id`.
-

## 12) Agentic Integration Hooks

- After report write: emit `report.created` (internal bus) and deliver webhooks.
  - Subscriber (agent orchestrator) clusters issues → proposes patches/tests → opens PR.
  - Optionally expose `/agents/actions` in future to attach `patch_unified` back to a job.
- 

## 13) Retention & Privacy (defaults)

- **Workspace** purged after 7 days; **reports** retained 30 days (configurable)
  - No PII processed unless present in repo; avoid storing access tokens in logs
  - Provide `/jobs/{job_id}` **DELETE** to purge workspace & report early
- 

## 14) FAQ

- **Why 200 vs 202 on `/analyze`?** Small repos may complete within request timeout → `200`. Larger ones return `202` and continue in background.
- **Do you execute user code?** No. Only static analysis tools run.
- **Can I scan private repos?** Yes via deploy token/credential injection (future); start with public repos.
- **How do I tune false positives?** Update `/config/analyzers` rulesets and per-tool ignore config; honor `.semgrepignore` and `# nosec` pragmas where feasible.